


MERIDIAN

Financial Research and News Intelligence

Project Case Study | October 2026

A connected workflow for inspecting company financials, valuation assumptions, earnings quality and news evidence.



Project overview

Meridian brings financial analysis and event context into a unified research interface. Its strongest demonstrated feature is the path from a summary assessment to the supporting detail: financial comparisons, valuation context, component scores, filing records and company relationships.

This case study explains the problem, walks through the demonstrated product and identifies concrete validation priorities. The screenshots document the project as captured on 4 October 2026.

Executive summary

Meridian is a financial research and news intelligence project that brings company analysis and event context into a connected interface. The demonstrated company workspace includes an overview, three financial statements, valuation, historical trends, earnings quality, financial health, SEC filings and relevant events. A separate news workspace combines a globe, an event feed and an evidence panel.


The central product idea is to help a reviewer move from a conclusion to the information needed to question it. A research score is accompanied by component scores; valuation has a range and confidence label; earnings quality has evidence controls; and news relevance is distinguished from a simple company mention.

This report documents the product through 16 screenshots captured on 4 October 2026, focusing on the visible workflow, design decisions and reasoning behind the system.

My contribution

I designed and developed Meridian as an end-to-end research project, covering both the financial research workflow and the News Intelligence system. My work included building the data pipeline, organising financial statement and market data, developing the research scoring and valuation logic, and creating separate modules for historical trends, earnings quality, financial health and SEC filing review. I also worked on the News Intelligence pipeline, where news events are classified and linked to companies based on their actual relationships and exposures rather than simple keyword mentions. On the product side, I designed the interface around how I would personally review a company: starting from a high-level assessment, then moving into the underlying financial data, model assumptions and source evidence when something needs further investigation. The project is still being refined, particularly around model validation, data consistency and reproducibility, but the main goal has remained the same throughout: to build a research tool that makes it easier to understand not only what a model is saying, but also why it is saying it.

Meridian Research Workflow



From raw company, market and filing data to structured research outputs and evidence review.

Meridian News Intelligence Workflow



From raw news items to structured event interpretation and company-level relevance.

Problem and intended users

A company assessment can require several different kinds of evidence: financial statement changes, cash generation, market performance, valuation assumptions, source filings and recent events. When these are disconnected, a reviewer has to reconstruct the context repeatedly and may mistake a summary score for a complete explanation.

Meridian is organised around that review process. Its intended users include company researchers, financial analysts and people evaluating market information who need a structured starting point for further investigation.

The product objective is to make related evidence easier to inspect within one workflow. Reduced research time and improved decision quality are potential benefits to test, not measured outcomes established by these screenshots

Demonstrated scope

Area	Visible outcome	Purpose in the workflow
Company overview	Description, quote, research score, component scores and key metrics	Establish the initial research context
Financial statements	Income statement, cash flow and balance sheet views with comparisons	Identify changes requiring explanation
Valuation	Low, composite and high values; base DCF; multiples; confidence	Put a model estimate in context
Historical trends	Common-base price chart and six selectable financial metrics	Compare market behaviour with business development
Earnings quality	Six component scores and evidence controls	Break an overall assessment into reviewable parts
Financial health	Cash, debt, net cash and ratio context	Examine the balance sheet alongside profitability
Filing records	SEC forms, dates and accession identifiers	Provide a starting point for source verification
Company events	Event cards, relationship status and evidence labels	Separate a mention from a supported relationship
News Intelligence	Globe, feed, event details, source links and status indicators	Review events in geographical and company context

The screenshots show one company in detail. The company selector is visible, but the coverage and completeness of other companies are outside this report's evidence.

Research walkthrough

Start with the summary and identify a tension

The Apple overview displays a research score of 51.9/100 alongside Fundamental, Valuation and Momentum components. Earnings quality is separately displayed as 86.7/100 elsewhere in the workspace.

The useful design choice is that these dimensions are visible separately. A business can have apparently strong earnings characteristics while the model evaluates its valuation unfavourably. A reviewer can investigate the disagreement rather than assuming every signal must point in the same direction.




Figure 1. The captured overview separates the headline rating from its displayed components. The score formula is not established by this screenshot.

Compare the statements before accepting the summary

The financial workspace places annual comparisons beside historical charts. In the captured income statement, revenue rises from \$391.04B to \$416.16B and net income from \$93.74B to \$112.01B. In the cash-flow view, operating cash flow declines from \$118.25B to \$111.48B and free cash flow from \$108.81B to \$98.77B.

Taken as displayed, these values give the reviewer a concrete question: why do reported earnings improve while cash generation weakens? The screenshots alone do not establish the answer. The next research step would be to inspect tax, working-capital and other reconciling items for the same periods.

The interface also displays “Review required” and a “Deterministic evidence fallback” summary. That state is useful because it makes an unresolved interpretation visible. The fallback label does not, by itself, prove how the summary is generated or whether an AI model is involved.




Figure 2. The cash-flow view provides evidence for a follow-up question instead of relying only on the headline earnings assessment.

Inspect the valuation and its confidence

The valuation view shows a displayed quote of \$333.69, composite fair value of \$174.95, base DCF of \$109.55, and a range of \$79.11–\$268.11. It states that the composite combines relative and intrinsic valuation families at 50%/50% weights, and labels valuation confidence “Low.”

These are captured model outputs, not independently validated price targets. The methodology panel is collapsed, so the discount rate, forecast assumptions, terminal value and peer selection cannot be established from the image.

The interface labels one result “Expected return.” Using the rounded displayed numbers, $\$174.95 / \$333.69 - 1$ is approximately -47.6% , matching that label numerically. This is an arithmetic consistency check only. Without a horizon, probabilities or treatment of distributions, “gap to model fair value” would communicate the visible calculation more precisely.




Figure 3. Range and confidence are presented together. Reproducing the model requires inputs and methodology beyond the screenshot.

Examine history and the quality of earnings

Historical Trends displays AAPL, SPY and QQQ on a base-100 chart. It explicitly describes the comparison as price performance over available common history and states that it is not risk-adjusted alpha. The separate financial-metric selector includes revenue growth, EPS growth, operating margin, free-cash-flow margin, ROE and ROIC.

This supports two different questions: how the market series compare, and how financial indicators develop over time. The displayed difference versus SPY is expressed in percentage points. Price adjustment, dividends and missing-date handling still require implementation evidence before the chart can support a verified performance claim.




Figure 4. The interface distinguishes a market comparison from the annual financial metrics shown below it.

The earnings-quality screen separates Cash Conversion, Accrual Quality, Margin Stability, ROIC Quality, Non-Core Earnings and Revenue Resilience. Each has an evidence control. This organisation creates a useful review structure, but neither the component rules nor the aggregation weights are visible. The overall score should not be described as audited or predictive on the basis of the interface alone.




Figure 5. Component-level presentation offers a route to explanation; the scoring methodology needs separate documentation.

Follow the source records and related events

The SEC timeline lists three filing records with form types, dates and accession numbers. Its source-coverage note explicitly separates filing availability from evidence assessment. This distinction matters: having a filing record is not the same as having verified the financial claim attributed to it.




Figure 6. Source identifiers make later verification possible. The screenshot does not establish that every metric is linked to a specific filing passage.

The company-events page shows an AWS outage item as an Apple “Review Candidate” with “Mention only” evidence. The news workspace distinguishes Amazon's direct relationship from Apple's weaker association. This is an important product distinction: an entity mention is not automatically a material exposure, and event severity is not a forecast of share-price direction.




Figure 7. News is presented with source, classification, company relationship and confirmation context. The captured workspace is marked DEMO.

Design decisions visible in the result

Organise navigation around research questions. The section sequence moves from summary to statements, valuation, historical context, quality, financial health and evidence. That gives a reviewer a recognisable route through a company assessment.

Expose uncertainty next to the result. Low valuation confidence, review-required summaries, mention-only relationships and stale-data indicators make uncertainty visible where a user could otherwise overinterpret a score or headline.

Use progressive detail. Evidence and methodology controls keep the main view readable while indicating where deeper explanation belongs. The screenshots demonstrate the controls, not the completeness of their expanded content.

Keep different judgments separate. Earnings quality, valuation, market momentum, event severity and company relevance answer different questions. Their separation is one of the strongest aspects of the demonstrated product structure.

Use distinct workspaces for different tasks. The lighter company-research view emphasises tables and financial charts; the darker news workspace emphasises spatial context and event selection. Cross-links indicate how a company investigation can continue into news review.

Known limitations and validation priorities

The screenshots are evidence of the interface and captured outputs. They cannot verify data ingestion, end-to-end interaction, model accuracy, access control, reliability or deployment readiness. The word “production” appears in several screen labels; it is not treated here as proof of operational maturity.

The captured demo also provides concrete improvement priorities:

Observation	Why it matters	Proposed verification or improvement
The company-events card shows Medium severity for the AWS item, while News Intelligence shows N/A and not confirmed	A reader may receive different interpretations of the same item	Compare record identifiers, versions and rendering rules; explain or remove unintended differences
Quote, filing and financial periods appear in different contexts	A valuation can combine inputs from different dates	Display input dates together and document refresh policies
Scores are visible but their expanded methodology is not supplied	A precise score can imply more certainty than the evidence supports	Publish definitions, weights, thresholds and missing-data behaviour
Filing records are listed without visible claim-level excerpts	Source availability alone does not establish a claim	Connect important metrics and explanations to exact source locations
“Expected return” appears to match the fair-value gap	The wording may imply a forecast with a defined horizon	Define the metric and horizon or rename it to match the calculation
ROIC history uses September 30 labels while other views show fiscal-year-end dates	Normalised display dates may be mistaken for original observation dates	Document the time convention and preserve source dates
The news view has demo, stale and frozen labels	Users need to know whether the workspace is current	Keep these states explicit and test their propagation across views

These are review priorities, not claims that fixes or tests have already been completed. No numerical accuracy, latency, time-saving or investment-performance result is asserted.

What the project demonstrates

The artefact gives a reviewer concrete material for discussing research requirements, financial presentation, information hierarchy and uncertainty. For a business or data analyst role, the statement comparisons and evidence distinctions are especially relevant. For a financial research role, the link between cash generation, valuation assumptions and source verification provides a useful discussion path. For a product role, the connected navigation and review states show how a complex task can be broken into manageable decisions.

Next development milestones

1. **Document reproducibility:** publish metric definitions, formula inputs, period conventions and a traceable worked example.
2. **Align shared event states:** verify company-event and global-news views against the same event and relationship records.
3. **Validate calculation behaviour:** check a representative set of statements, ratios, valuation cases and missing-data cases against expected results.
4. **Evaluate the workflow with users:** measure completion time, source-verification success and interpretation errors on defined research tasks.

The next release should be assessed against these observable outcomes rather than the number of additional dashboard features.

Evidence and scope

Figures reproduce project screenshots captured on 4 October 2026. Displayed prices, ratings, headlines and model outputs illustrate the application at that point in development and should not be treated as current market information or investment recommendations. This case study focuses on the demonstrated product workflow and visible outputs; further validation of model behaviour, data consistency and reproducibility remains part of the ongoing development process.